

# 1 Analyzer

Can you explain the performance differences between the different versions using some analysis tool, e.g. analyzer? Does the tool confirm your expectations?

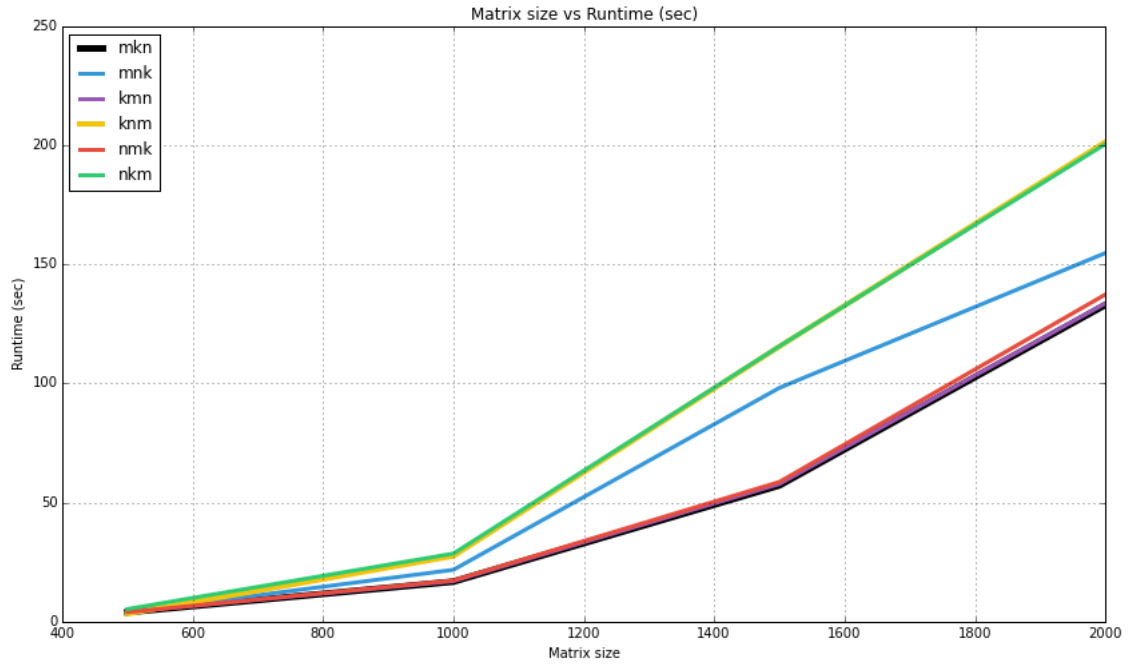


Figure 1: Plot for the six different permutations, matrix sizes vs. runtime, square matrices of 500, 1000, 1500 and 2000.

Permutations	m=n=k=500	m=n=k=1000	m=n=k=1500	m=n=k=2000
mkn	4.050 s	16.880 s	57.250 s	132.990 s
kmn	4.110 s	17.070 s	57.820 s	133.580 s
nmk	4.080 s	17.370 s	58.710 s	137.330 s
mnk	4.080 s	21.840 s	98.160 s	154.730 s
nkm	5.380 s	28.650 s	115.640 s	200.460 s
knm	3.410 s	27.610 s	115.520 s	201.300 s

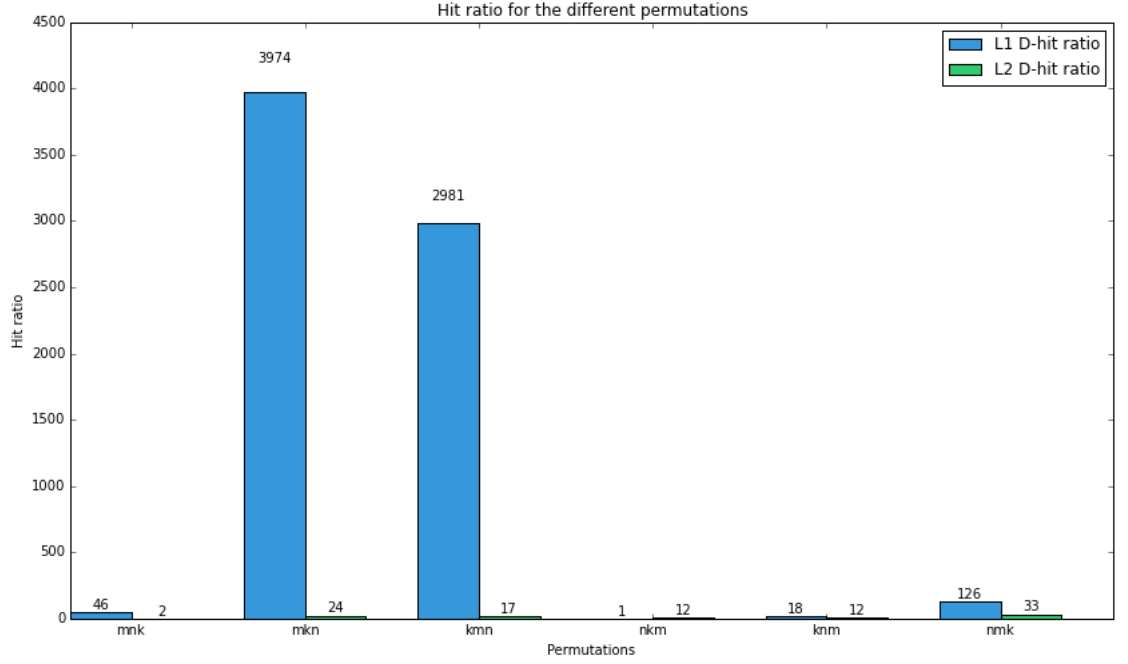


Figure 2: Plot for hit ratio vs. permutations, square matrix of 500, mflops\_max\_it of 50.

$$CacheRefs - CacheMisses / CacheMisses = HitRatio$$

L2 D	mnk	mkn	kmn	nkm	knm	nmk
Cache Refs	100815940823	150726065855	150738548844	15124047909	151263541727	150757006506
Cache Misses	2169988937	37921061	50562333	8271302759	8254503877	1192497865
Hit Ratio	46.4591	3974.7322	2981.2419	1.8284	18.3249	126.4211

L1 D	mnk	mkn	kmn	nkm	knm	nmk
Cache Refs	2337185157	808373113	816761066	9726740774	9827651083	2025085625
Cache Misses	790190116	33302649	46629262	808869156	794451254	59968558
Hit Ratio	2.9597	24.27	17.516	12.025	12.3703	33.7691

A cache miss, generally, is when something is looked up in the cache and is not found. The cache did not contain the item being looked up. The cache hit is when you look something up in a cache and it was storing the item and is able to satisfy the query.

We saw previously that "mkn" has the lowest runtime followed by "kmn". Therefore, they should have the highest hit ratio. We can clearly see in figure 2 that the expectation is confirmed. The faster the algorithm runs the higher hit ratio it has. This makes sense. In order to have a high hit

ratio, there are many hits and few misses. If there are many hits the data is stored in the local memory cache. Therefore, the thread runs fast.

The slower the algorithm runs the lower is the hit ratio. This is understandable as well. In order to have a low hit ratio, there are few hits and many misses. If there are many misses, this means the data is not stored in the local memory cache. The hardware has to then make a number of requests to main memory to fill up the local memory cache which causes the thread to run slower.